

Convolution and Hough Transform

March 11, 2022

1 Lab Assignment 2

Computer Vision - Term 5, 2022

Instructor: Dr. Saumya Jetly TA: Ribhu Lahiri

Deadline: Sunday, 13 March 2022 11:59 am

Submission form link: <https://forms.gle/HGkVEoMgK62C7oWd7>

Total points: 5 (with possible extra credit)

1.0.1 Task 1: Creating and applying new filters (3 points)

The first task is to create the convolve function. As discussed in lecture, you need to implement a function which takes a filter (kernel) and convolves it over the image using a sliding window. As an output you should get the processed image.

Extra Credit: Create a mathematical convolution function and a correlation function. Use both on the same image with the same filter. Is there a difference in the output? (0.5 points)

```
[3]: # Imports
import numpy as np
import pandas as pd
import cv2
import matplotlib.pyplot as plt
import scipy.fftpack as fp
```

```
[ ]: # Reading in a sample image
from PIL import Image

im = cv2.imread("lena.png", 0)
```

Implement the convolve function (2 points)

Use the helper method `sliding_window` to create a function that uses the sliding window to convolve over a given image.

```
[ ]: # Task:

def sliding_window(receptive_field, filter):
    '''
```

```

Convolves a filter over the receptive field from the image

Parameters
-----
receptive_field: np.ndarray
    The portion of the image the convolution is being done over

filter: np.ndarray
    The defined filter

Returns
-----
np.ndarray
    The convolved receptive field output

'''

return np.sum(filter * receptive_field)

def convolve(image, kernel):
    '''
    Convolves the filter over the image (sliding window)

    Parameters
    -----
    image: np.ndarray
        The image as a 2-dimensional matrix

    kernel: np.ndarray
        The defined filter

    Returns
    -----
    output_image: np.ndarray
        The image after convolving the filter over it
    '''

    image_size = image.shape
    kernel_size = kernel.shape
    # Zero padding to return output of same size as input
    if image_size[0]%2!=0:
        image=cv2.resize(image,(image.shape[0]-1,image.shape[1]))
    if image_size[1]%2!=0:
        image=cv2.resize(image,(image.shape[0],image.shape[1]-1))

```

```

    # Calculating the number of zeros to be padded to get the output size same
    ↪ as input
    x_padding_size=((image_size[0]+kernel_size[0]-1)-image_size[0])//2
    y_padding_size=((image_size[1]+kernel_size[1]-1)-image_size[1])//2

    # Zero Padding
    image=np.vstack((np.zeros((y_padding_size,image.shape[1])),image))
    image=np.vstack((image,np.zeros((y_padding_size,image.shape[1]))))
    image=np.hstack((image,np.zeros((image.shape[0],x_padding_size))))
    image=np.hstack((np.zeros((image.shape[0],x_padding_size)),image))

    #Flipping Kernel to do convolution
    kernel = np.flipud(np.fliplr(kernel))

    image_size = image.shape[0]
    kernel_size = kernel.shape[0]

    # Creating two for loops to iterate over different sliding windows
    convoluted_image=[]
    for x_slide in range(0,image_size-kernel_size+1):
        temp_result=[]
        for y_slide in range(0,image_size-kernel_size+1):
            # Now we have to get the sub image to pass to the calculator function
            window=[]
            for i in range(x_slide,x_slide+kernel_size):
                temp=[]
                for j in range(y_slide,y_slide+kernel_size):
                    temp.append(image[i][j])
                window.append(temp)
            temp_result.append(int(sliding_window(np.array(window),np.
    ↪ array(kernel)).sum()))
            convoluted_image.append(temp_result)
        return(np.array(convoluted_image))

def correlate(image, kernel):
    """
    Convolves the filter over the image (sliding window)

    Parameters
    -----
    image: np.ndarray
        The image as a 2-dimensional matrix

```

```

kernel: np.ndarray
    The defined filter

Returns
-----
output_image: np.ndarray
    The image after convolving the filter over it
'''
image_size = image.shape
kernel_size = kernel.shape
    # Zero padding to return output of same size as input
if image_size[0]%2!=0:
    image=cv2.resize(image,(image.shape[0]-1,image.shape[1]))
if image_size[1]%2!=0:
    image=cv2.resize(image,(image.shape[0],image.shape[1]-1))

    # Calculating the number of zeros to be padded to get the output size same
→as input
x_padding_size=((image_size[0]+kernel_size[0]-1)-image_size[0])//2
y_padding_size=((image_size[1]+kernel_size[1]-1)-image_size[1])//2

# Zero Padding
image=np.vstack((np.zeros((y_padding_size,image.shape[1])),image))
image=np.vstack((image,np.zeros((y_padding_size,image.shape[1]))))
image=np.hstack((image,np.zeros((image.shape[0],x_padding_size))))
image=np.hstack((np.zeros((image.shape[0],x_padding_size)),image))

image_size = image.shape[0]
kernel_size = kernel.shape[0]

# Creating two for loops to iterate over different sliding windows
convoluted_image=[]
for x_slide in range(0,image_size-kernel_size+1):
    temp_result=[]
    for y_slide in range(0,image_size-kernel_size+1):
        # Now we have to get the sub image to pass to the calculator function
        window=[]
        for i in range(x_slide,x_slide+kernel_size):
            temp=[]
            for j in range(y_slide,y_slide+kernel_size):
                temp.append(image[i][j])
            window.append(temp)
        temp_result.append(int(sliding_window(np.array(window),np.
→array(kernel)).sum()))
        convoluted_image.append(temp_result)
    return(np.array(convoluted_image))

```

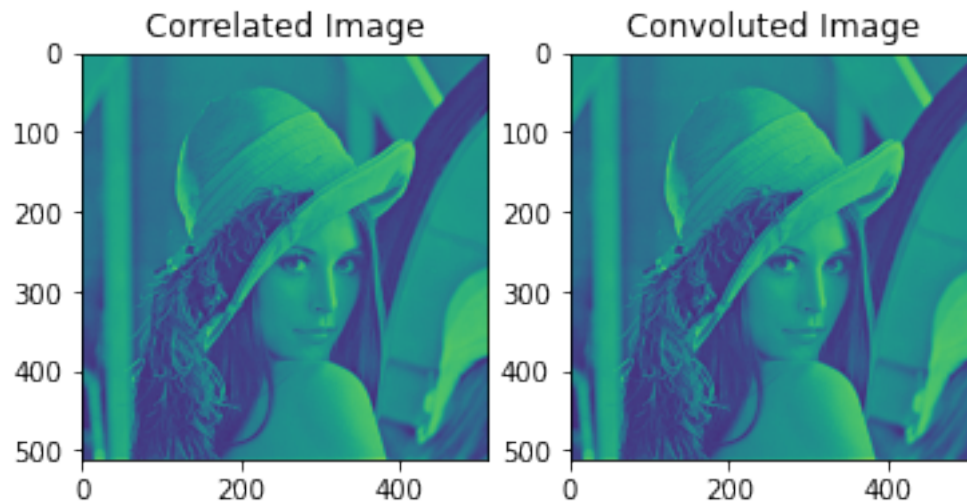
```
[ ]: # Applying convolution and Correlation to same image to check the output
```

```
im = np.array(im)
filter = np.array([[0, -1, 0],
                  [-1, 7, -1],
                  [0, -1, 0]])
correlated=correlate(im,filter)
convoluted=convolve(im,filter)

plt.subplot(1,2,1)
plt.imshow(correlated)
plt.title('Correlated Image')

plt.subplot(1,2,2)
plt.imshow(convoluted)
plt.title('Convolved Image')
```

```
[ ]: Text(0.5, 1.0, 'Convolved Image')
```



```
[ ]: correlated[correlated!=convoluted] # Returns None if all values are same
```

```
[ ]: array([], dtype=int64)
```

1.0.2 It can be seen that both from the images and values, the result of convolution and correlation are the same

Sample Sharpen filter to check functionality


```
[ ]: image_array = np.asarray(im)
```

```
# Sharpen
filter = np.array([[0, -1, 0],
                  [-1, 7, -1],
                  [0, -1, 0]])
```

```
[ ]: filtered_im = convolve(image_array, filter)
```

```
[ ]: fsize = (10,6)
plt.figure(figsize = fsize)
plt.subplot(121),plt.imshow(im, 'gray'), plt.title('Original')
plt.subplot(122),plt.imshow(filtered_im, 'gray'), plt.title('Sharpened')
```

```
[ ]: (<matplotlib.axes._subplots.AxesSubplot at 0x7feb94260150>,
     <matplotlib.image.AxesImage at 0x7feb96c49750>,
     Text(0.5, 1.0, 'Sharpened'))
```



Sobel filters to check functionality

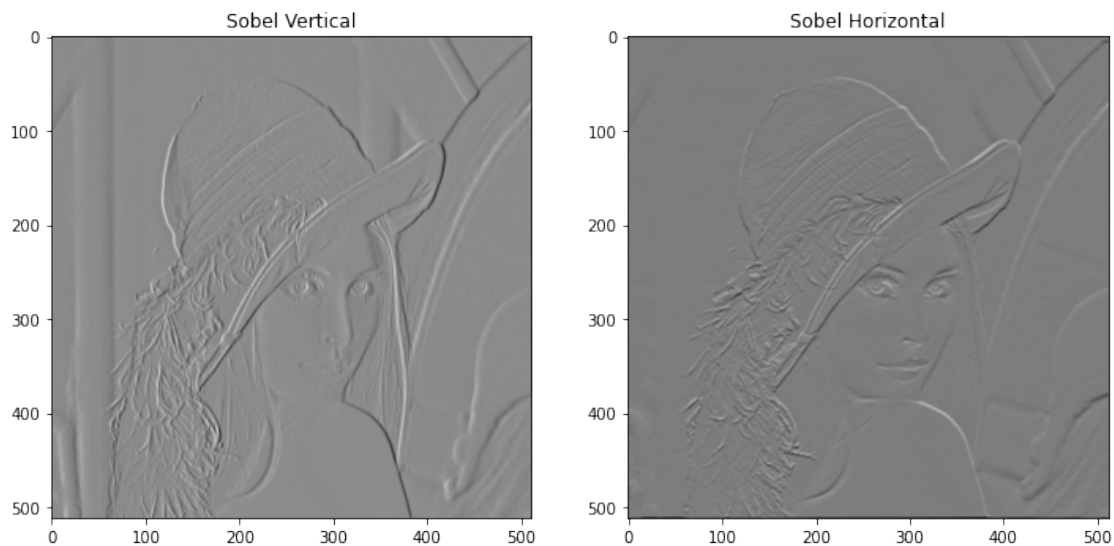
```
[ ]: # Sobel filters
sobel_vert = np.array([[1, 0, -1],
                      [2, 0, -2],
                      [1, 0, -1]])

sobel_hor = np.array([[1, 2, 1],
                      [0, 0, 0],
                      [-1, -2, -1]])
```

```
filtered_v = convolve(image_array, sobel_vert)
filtered_h = convolve(image_array, sobel_hor)
```

```
[ ]: fsize = (12,8)
plt.figure(figsize = fsize)
plt.subplot(121), plt.imshow(filtered_v, 'gray'), plt.title('Sobel Vertical')
plt.subplot(122), plt.imshow(filtered_h, 'gray'), plt.title('Sobel Horizontal')
```


```
[ ]: (<matplotlib.axes._subplots.AxesSubplot at 0x7feb9180cf90>,
<matplotlib.image.AxesImage at 0x7feb9372cc90>,
Text(0.5, 1.0, 'Sobel Horizontal'))
```



```
[ ]: #Combining them
sobel_edge_detector = np.sqrt(filtered_h**2 + filtered_v**2)
```

```
[ ]: plt.figure(figsize=fsize)
plt.imshow(sobel_edge_detector, 'gray')
```

```
[ ]: <matplotlib.image.AxesImage at 0x7feb9392bdd0>
```



Try the convolve function with your own filter (1 points)

Create your own filter (or use one from Google) and convolve it over your imported image


```
[ ]: filter = np.array([[0, 2, 0],  
                        [2, 0, 2],  
                        [0, 2, 0]])
```

```
[ ]: filtered_im = convolve(image_array, filter)
```

```
[ ]: fsize = (10,6)  
plt.figure(figsize = fsize)  
plt.subplot(121),plt.imshow(im, 'gray'), plt.title('Original')  
plt.subplot(122),plt.imshow(filtered_im, 'gray'), plt.title('Filtered')
```



```
[ ]: (<matplotlib.axes._subplots.AxesSubplot at 0x7feb915a7c90>,
<matplotlib.image.AxesImage at 0x7feb9145d510>,
Text(0.5, 1.0, 'Filtered'))
```



1.0.3 Task 2: Creating and applying Hough filter (2 points)

Implement Hough Transform (2 points)

Create a Hough Transform to detect **rectangles** in a given image. It might not be perfect due to the hyperparameters you pick but your goal should be to implement the function.

Feel free to use the helper canny function along with any other inbuilt opencv methods

Extra Credit: Implement the Hough Transform from scratch, i.e. without using openCV (0.5 points)

```
[180]: def canny(img):
        """
        Applies the canny filter to detect edges in a given image

        Parameters
        -----
        img: np.ndarray
            The image as a 2-dimensional matrix

        Returns
        -----
        edges: np.ndarray
            The image after applying canny filter, i.e. with edges
            highlighted
```

```

'''
low_threshold = 50
high_threshold = 800
edges = cv2.Canny(img, low_threshold, high_threshold)
return edges

def hough(img):
'''
    Apply Hough Transform to a given image to detect rectangles.

    Parameters
    -----
    img: np.ndarray
        The image as a 2-dimensional matrix

    Returns
    -----
    edges: np.ndarray
        The image after applying hough transform, i.e. with rectangles
        highlighted
'''
edges = canny(img)
lines = cv2.HoughLinesP(edges,1,np.pi/180,5,1000)
for line in lines:
    x1,y1,x2,y2=line[0]
    cv2.line(img,(x1,y1),(x2,y2),(255,0,0),2)
return img

```

```


[181]: im = cv2.imread("box.jpg")
plt.imshow(hough(im))

```

```

[181]: <matplotlib.image.AxesImage at 0x7f55b0773650>

```



Manually done Hough transform code attached below. The execution was take a large amount of time and hence could not implement.

```
[170]: im = cv2.imread("box.jpg")
# The function returns the pixels at which the edges are detected
# Iterating through the result to get details about the coordinates which are
# identified in Canny edge detection
import pandas as pd
pixels=np.array(canny(im))
coordinates=[]
count=0
for i in range(pixels.shape[0]):
    for j in range(pixels.shape[1]):
        if pixels[i][j]>0:
            coordinates.append((i,j))
        else:
            count+=1

# Creating an accumulator for  $\theta$  and  $d$ 
accumulator_size = 20

accumulator=np.zeros((accumulator_size,accumulator_size))
theta = np.linspace(0,180,accumulator_size)
x_variable = [coordinate[0] for coordinate in coordinates]
y_variable = [coordinate[1] for coordinate in coordinates]
```

```

# Creating the equation for filling values inside the accumulator
from sympy import *
from tqdm import tqdm

x, y, t= symbols(' x y t')
polar_expression = (x*cos(t))+ (y*sin(t))

# Defining the bins for equation results
binInterval = np.linspace(-50,50,20)
binLabels    = np.arange(0,19)
# Solving the equation for each pixel values for different values of theta
results=[]
for i in tqdm(x_variable):
    for j in y_variable:
        for k in theta:
            results.append(float((i*np.cos(k))+(j*np.sin(k))))
df=pd.DataFrame()
df['data']=results
df['binned'] = pd.cut(df['data'], bins = binInterval, labels=binLabels)
bins=df['binned'].values
for bin_number in range(len(theta)):
    try:
        accumulator[bins[bin_number]][bin_number]+=1
    except:
        print(bin_number)

```

```

0%|          | 0/3040 [00:00<?, ?it/s]
0%|          | 1/3040 [01:34<79:42:17, 94.42s/it]

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-170-ff156b724a23> in <module>()
    40     results.append(float((i*np.cos(k))+(j*np.sin(k))))
    41     df=pd.DataFrame()
--> 42     df['data']=results
    43     df['binned'] = pd.cut(df['data'], bins = binInterval,
    ↪ labels=binLabels)
    44     bins=df['binned'].values

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py in __setitem__(self,
    ↪ key, value)
    3610         else:
    3611             # set column
-> 3612             self._set_item(key, value)
    3613

```

```

3614     def _setitem_slice(self, key: slice, value):

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py in _set_item(self,
↳key, value)
    3782         ensure_homogeneity.
    3783         """
-> 3784         value = self._sanitize_column(value)
    3785
    3786         if (

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py in
↳_sanitize_column(self, value)
    4500         numpy.ndarray or ExtensionArray
    4501         """
-> 4502         self._ensure_valid_index(value)
    4503
    4504         # We should never get here with DataFrame value

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py in
↳_ensure_valid_index(self, value)
    3856             index_copy.name = self.index.name
    3857
-> 3858             self._mgr = self._mgr.reindex_axis(index_copy, axis=1,
↳fill_value=np.nan)
    3859
    3860     def _box_col_values(self, values, loc: int) -> Series:

/usr/local/lib/python3.7/dist-packages/pandas/core/internals/base.py in
↳reindex_axis(self, new_index, axis, fill_value, consolidate, only_slice)
     85         Conform data manager to new index.
     86         """
---> 87         new_index, indexer = self.axes[axis].reindex(new_index)
     88
     89         return self.reindex_indexer(

/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in
↳reindex(self, target, method, level, limit, tolerance)
    3824             if self._index_as_unique:
    3825                 indexer = self.get_indexer(
-> 3826                     target, method=method, limit=limit,
↳tolerance=tolerance
    3827                 )
    3828             else:

/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in
↳get_indexer(self, target, method, limit, tolerance)
    3484         )
    3485

```

```

-> 3486         return self._get_indexer(target, method, limit, tolerance)
      3487
      3488     def _get_indexer(

/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in _
-> _get_indexer(self, target, method, limit, tolerance)
      3502         target = target.astype(dtype, copy=False)
      3503         return this.get_indexer(
-> 3504             target, method=method, limit=limit, tolerance=tolerance
      3505         )
      3506

/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in _
-> get_indexer(self, target, method, limit, tolerance)
      3484         )
      3485
-> 3486         return self._get_indexer(target, method, limit, tolerance)
      3487
      3488     def _get_indexer(

/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in _
-> _get_indexer(self, target, method, limit, tolerance)
      3510         indexer = self._get_nearest_indexer(target, limit, tolerance)
      3511     else:
-> 3512         indexer = self._engine.get_indexer(target.
-> _get_engine_target())
      3513
      3514         return ensure_platform_int(indexer)

```

KeyboardInterrupt: